

1 Substrings

- 1.1 You wish to speed up one of your programs by implementing it directly in assembly. Your partner started translating the function `is_substr()` from C to RISC-V, but didn't finish. Please complete the translation by filling in the lines below with RISC-V assembly. The prologue and epilogue have been written correctly but are not shown.

Note: `strlen()`, both as a C function and RISC-V procedure, takes in one string as an argument and returns the length of the string (not including the null terminator).

```
/* Returns 1 if s2 is a substring of s1, and 0 otherwise. */
```

```
int is_substr(char* s1, char* s2) {  
    int len1 = strlen(s1);  
    int len2 = strlen(s2);  
    int offset = len1 - len2;  
    while (offset >= 0) {  
        int i = 0;  
        while (s1[i + offset] == s2[i]) {  
            i += 1;  
            if (s2[i] == '\0')  
                return 1;  
        }  
        offset -= 1;  
    }  
    return 0;  
}
```

1.2 Fill in the following RISC-V code based on the given C code:

```

1. is_substr:
2.  mv s1, a0
3.  mv s2, a1
4.  jal ra, strlen
5.  mv s3, a0
6.  mv a0, s2
7.  jal ra, strlen
8.  sub s3, s3, a0
9. Outer_Loop:
10. _____, _____, _____, False
11. add t0, x0, x0
12. Inner_Loop:
13. add t1, t0, s3
14. add t1, s1, t1
15. lbu t1, 0(t1)
16. _____
17. _____
18. _____, t1, _____, Update_Offset
19. addi t0, t0, 1
20. add t2, t0, s2
21. _____
22. beq t2, _____, _____,
23. jal x0 Inner_Loop
24. Update_Offset: addi s3, s3, -1
25. _____
26. False: xor a0, a0, _____
27. jal x0, End
28. True: addi a0, x0, 1
29. End: _____.
```

2 Advanced RISC-V

2.1 You are given the following RISC-V code:

```
Loop:   andi t2 t1 1
        srli t3 t1 1
        bltu t1 a0 Loop
        jalr s0 s1 MAX_POS_IMM
```

- (a) What is the value of the byte offset that would be stored in the immediate field of the bltu instruction?
- (b) What is the binary encoding of the bltu instruction? Please use hexadecimal to represent your answer.

2.2 As a curious 61C student, you question why there are so many possible opcodes, but only 47 instructions. Thus, you propose a revision to the standard 32-bit RISC-V instruction formats where each instruction has a unique opcode (which still is 7 bits). You believe this justifies taking out the funct3 field from the R, I, S, and SB instructions, allowing you to allocate bits to other instruction fields except the opcode field.

- (a) What is the largest number of registers that can now be supported in hardware?
- (b) With the new register size, how far can a jal instruction jump to (in halfwords)?
- (c) Assume register `s0 = 0x1000 0000`, `s1 = 0x4000 0000`, `PC = 0xA000 0000`. Lets analyze the instruction `jalr s0, s1, MAX_POS_IMM` where `MAX_POS_IMM` is the maximum possible positive immediate for `jalr`. Using the register sizes defined above, what are the values in registers `s0`, `s1`, and `pc` after the instruction executes?

3 Cache Rules Everything Around Me

- 3.1 You are given a RISC-V machine with a single level of 2KiB direct-mapped cache with 512B cache blocks. It has 1MiB of physical address space.

The function `foo()` is ran on the system with a cold cache and as the only process:

```
#define ARRAY_LEN 4096
#define STEP_SIZE 64
// A starts at 0x10000
// B starts at 0x20000
void foo (int* A, int* B) {
    int total = 0;
    for ( int i = 0; i < ARRAY_LEN; i += STEP_SIZE )
    {
        total += A[ i ];
        total -= B[ i ];
    }
}
```

- (a) Calculate the number of Tag, Index, and Offset bits for this cache.

- (b) Calculate the hit percentage for this cache after running `foo`.

- (c) The cache is now cleared and the code is run again. This time, A and B are pointing to the same array, which starts at `0x10000`. Calculate the new hit percentage.

- (d) Assume A and B starts once again at `0x10000` and `0x20000`. What is the new hit percentage if we ran `foo` on a fully associative cache, with all other parameters staying the same?