

1 Data Level Parallelism

<code>__m128i _mm_set1_epi32(int i)</code>	sets the four signed 32-bit integers to i
<code>__m128i _mm_loadu_si128(__m128i *p)</code>	returns 128-bit vector stored at pointer p
<code>__m128i _mm_add_epi32(__m128i a, __m128i b)</code>	returns vector (a0+b0, a1+b1, a2+b2, a3+b3)
<code>__m128i _mm_mullo_epi32(__m128i a, __m128i b)</code>	returns vector (a0*b0, a1*b1, a2*b2, a3*b3)
<code>__m128i _mm_cmpgt_epi32(__m128i a, __m128i b)</code>	returns: vector((a0>b0)?0xffffffff:0, (a1>b1)?0xffffffff:0, (a2>b2)?0xffffffff:0, (a3>b3)?0xffffffff:0)
<code>void _mm_storeu_si128(__m128i *p, __m128i a)</code>	stores 128-bit vector a at pointer p

1.1 The following code uses loop unrolling to improve performance:

```
static void sum_unrolled (int *a) {
    int sum = 0;
    for (int i = 0; i < 16; i += 4) {
        sum += (a + i);
        sum += (a + i + 1);
        sum += (a + i + 2);
        sum += (a + i + 3);
    }
}
```

(a) Implement the following function with a single SIMD instruction:

```
static void sum_unrolled_SIMD (int *a) {
    int sum = 0;
    int result[4];
    __m128i result_v = _mm_set1_epi32(0);
    for (int i = 0; i < 16; i += 4) {
        result_v = _____;
    }
    _mm_storeu_si128((__m128i*)result, result_v);
    sum = result[0] + result[1] + result[2] + result[3];
}
```

1.2 (a) Implement the following function using SIMD:

```
// Sequential code
static int selective_sum_total (int n, int *a, int c) {
    int sum = 0;
    for (int i = 0; i < n; i += 1) {
        if (a[i] > c) {
            sum += a[i];
        }
    }
    return sum;
}

// SIMD code
static int selective_sum_vectorized (int n, int *a, int c) {
    int result[4];
    _m128i sum_v = _____;
    _m128i cond_v = _____;

    for (int i = 0; i < ____; i += ____) { //Vectorized loop
        _m128i curr_v = _mm_loadu_si128(_____);
        __m128i tmp = _mm_cmpgt_epi32(_____);
        sum_v = _____;
    }
    _mm_storeu_si128(_____);

    for (int i = ____; i < ____; i += 1) { //Tail case
        result[0] += _____;
    }
    return _____;
}
```

2 Thread Level Parallelism

Some OpenMP syntax:

<code>#pragma omp parallel{ ... }</code>	Signals the system to spawn threads
<code>#pragma omp for</code>	Split the for loop into equal-sized chunks
<code>int omp_get_num_threads(void);</code>	Returns the number of threads the system has
<code>int omp_get_thread_num(void);</code>	Returns the thread ID of the current thread

2.1 Implement the following function using openMP:

// Sequential code

```
static int selective_product_total (int n, int *a, int c) {
    for (int i = 0; i < n; i += 1) {
        a[i] = (a[i] > c)? 10 : 0;
    }
}
```

- (a) Given $\text{sizeof}(a^*) = \text{sizeof}(\text{int}) \cdot 16$, manually distribute the iterations such that there are four contiguous chunks and no false sharing. Assume the total number of threads is a multiple of 4. `int omp_get_num_threads` \leftarrow returns total number of threads in a team `int omp_get_thread_num` \leftarrow returns current thread ID number

// openMP code

```
static int selective_product_parallelized (int n, int *a, int c) {
    #pragma omp parallel {
        for (int i = 0; i < n; i += 1) {
            if ( _____ ) {
                a[i] = (a[i] > c)? 10 : 0;
            }
        }
    }
}
```

2.2 Implement the following function using openMP:

```
// Sequential code
static int selective_square_total (int n, int *a, int c) {
    for (int i = 0; i < n; i += 1) {
        if (a[i] > c) {
            a[i] *= a[i];
        }
    }
    return product;
}

(a) // openMP for code
static int selective_square_parallelized (int n, int *a, int c) {
    #pragma omp parallel {
        #pragma omp for
        for (int i = ____; i < ____; i += ____) {
            if (_____ > _____) {
                a[i] *= _____;
            }
        }
    }
    return product;
}
```

2.3 Assume the code is run with the following function call on a 32-bit machine and access to 8 threads. Also assume that ptr is block-aligned:

```
selective_square_parallelized(512, ptr, 61);
```

- (a) Assume the machine uses a 1 KiB direct-mapped cache with 256 B blocks. Is the code correct?

- (b) Assume the machine uses a 1 KiB direct-mapped cache with 1024 B blocks. Is the code correct?

3 Data Race!

- 3.1 Suppose we have `int *A` that points to the head of an array of length `len`. Determine which statement (A)-(E) correctly describes the code execution and provide a one or two sentence justification.

(a) Consider the following code:

```
#pragma omp parallel for
for (int x = 0; x < len; x++){
    *A = x;
    A++;
}
```

Is the code:

- A) Always Incorrect
- B) Sometimes Incorrect
- C) Always Correct, Slower than Serial
- D) Always Correct, Speed relative to Serial depends on Caching Scheme
- E) Always Correct, Faster than Serial

(a) `#pragma omp parallel`
`{`
 `for (int x = 0; x < len; x++){`
 `*(A+x) = x;`
 `}`
`}`

Is the code:

- A) Always Incorrect
- B) Sometimes Incorrect
- C) Always Correct, Slower than Serial
- D) Always Correct, Speed relative to Serial depends on Caching Scheme
- E) Always Correct, Faster than Serial