1 Data Level Parallelism

__m128i _mm_set1_epi32(int i)	sets the four signed 32-bit integers to i
__m128i _mm_loadu_si128(__m128i *p)	returns 128-bit vector stored at pointer p
__m128i _mm_add_epi32(__m128i a, __m128 b)	returns vector (a0+b0, a1+b1, a2+b2, a3+b3)
__m128i _mm_mullo_epi32 (__m128 a, __m128 b)	returns vector (a0*b0, a1*b1, a2*b2, a3*b3)
__m128i _mm_cmpgt_epi32(__m128 a, __m128 b)	returns: vector((a0>b0)?0xffffffff:0, (a1>b1)?0xffffffff:0, (a2>b2)?0xffffffff:0, (a3>b3)?0xffffffff:0)
void _mm_storeu_si128(__m128i *p,__m128i a)	stores 128-bit vector a at pointer p

1.1 The following code uses loop unrolling to improve performance:

```
static void sum_unrolled (int *a) {
    int sum = 0;
    for (int i = 0; i < 16; i += 4) {
        sum += (a + i);
        sum += (a + i + 1);
        sum += (a + i + 2);
        sum += (a + i + 3);
    }
}
```

(a) Implement the following function with a single SIMD instruction:

```
static void sum_unrolled_SIMD (int *a) {
    int sum = 0;
    int result[4];
    __m128i result_v = _mm_set1_epi32(0);
    for (int i = 0; i < 16; i += 4) {
        result_v = _____;
    }
    _mm_storeu_si128((__m128i*)result, result_v);
    sum = result[0] + result[1] + result[2] + result[3];
}
```

1.2   (a) Implement the following function using SIMD:

```
// Sequential code
static int selective_sum_total (int n, int *a, int c) {
    int sum = 0;
    for (int i = 0; i < n; i += 1) {
        if (a[i] > c) {
            sum += a[i];
        }
    return sum;
}


// SIMD code
static int selective_sum_vectorized (int n, int *a, int c) {
    int result[4];
    _m128i sum_v = _____;
    _m128i cond_v = _____;

    for (int i = 0; i < ___; i += ___) { //Vectorized loop
      _m128i curr_v = m128i_mm_loadu_si128(_____);
        __m128i tmp = _mm_cmpgt_epi32(_____);
        sum_v = _____;
    }
  _mm_storeu_si128(_____);

    for (int i = ___; i < ___; i += 1) { //Tail case
        result[0] += _____;
    }
    return _____;
```

1.3   (a) You have just begun your virtual internship at Machine Learning Inc.
Your boss (since you can't bring her coffee physically) wants you to run
data of her preferences through the company's coffee-classifying neural
network to decide which flavors she'll like best. You need to get it done
before next morning in order to impress her, but you notice that things
are running quite slowly - you dig deeper and find that the code's matrix
multiplications are not optimized for your machine!

Use SIMD instructions to rewrite the company's matrix-vector multi-
plication function so that it runs more efficiently on your computer.
The matrix has r rows and c columns. You may assume the length of
the vector is equal to c.

```c
// Company code
static int* matVecMul(int** mat, int* vec, int r, int c) {
    int sum;
    int* result = malloc(r * sizeof(int));
    for (int i = 0; i < r; i++) {
        sum = 0;
        for (int j = 0; j < c; j++) { sum += mat[i][j] * vec[j]; }
        result[i] = sum;
    }
    return result;
}


// SIMD code
static int* efficientMatVecMul(int** mat, int* vec, int r, int c) {
    __m128i sum_v;
    int* result = malloc(_____);
    for (int i = 0; i < r; _____) {
        sum_v = _mm_set1_epi32(0)
        for (int j = 0; j < _____; _____) {
            __m128i mat_v = _mm_loadu_si128(_____);
            __m128i vec_v = _mm_loadu_si128(_____);
            sum_v = _____;
        }
        int* sum_arr = malloc(_____);
        _mm_storeu_si128(sum_arr, sum_v);
        for (int k = _____; _____; _____) { sum_arr += _____; }
        result[i] = _____;
    }
    return result;
}
```

2   Thread Level Parallelism

Some OpenMP syntax:

`#pragma omp parallel{ ... }`	Signals the system to spawn threads
`#pragma omp for`	Split the for loop into equal-sized chunks
`int omp_get_num_threads(void);`	Returns the number of threads the system has
`int omp_get_thread_num(void);`	Returns the thread ID of the current thread

2.1   Implement the following function using openMP:

```
// Sequential code
static int selective_product_total (int n, int *a, int c) {
    for (int i = 0; i < n; i += 1) {
        a[i] = (a[i] > c)? 10 : 0;
}
```

(a) Given sizeof(a*) = sizeof(int)*16, manually distribute the iterations such that there are four contiguous chunks and no false sharing. Assume the total number of threads is a multiple of 4. int omp_get_num_threads ← returns total number of threads in a team int omp_get_thread_num ← returns current thread ID number

```
// openMP code
static int selective_product_parallelized (int n, int *a, int c) {
    #pragma omp parallel {
    for (int i = 0; i < n; i += 1) {
        if ( _____) {
                a[i] = (a[i] > c)? 10 : 0;
            }
        }
    }
}
```

2.2 Implement the following function using openMP:

```
// Sequential code
static int selective_square_total (int n, int *a, int c) {
    for (int i = 0; i < n; i += 1) {
        if (a[i] > c) {
            a[i] *= a[i];
        }
    }
    return product;
}
```

(a) ```
// openMP for code
static int selective_square_parallelized (int n, int *a, int c) {
    #pragma omp parallel {
    #pragma omp for
    for (int i = ___; i < ___; i += ___) {
        if (_____ > _____) {
            a[i] *= _____;
        }
    }
    return product;
}
```

2.3  Assume the code is run with the following function call on a 32-bit machine and access to 8 threads. Also assume that ptr is block-aligned:

```
selective_square_parallelized(512, ptr, 61);
```

(a) Assume the machine uses a 1 KiB direct-mapped cache with 256 B blocks. Is the code correct?

(b) Assume the machine uses a 1 KiB direct-mapped cache with 1024 B blocks. Is the code correct?

3 Data Race!

3.1 Suppose we have int *A that points to the head of an array of length len. Determine which statement (A)-(E) correctly describes the code execution and provide a one or two sentence justification.

(a) Consider the following code:

```
#pragma omp parallel for
for (int x = 0; x < len; x++){
    *A = x;
    A++;
}
```

Is the code:

A) Always Incorrect

B) Sometimes Incorrect

C) Always Correct, Slower than Serial

D) Always Correct, Speed relative to Serial depends on Caching Scheme

E) Always Correct, Faster than Serial

(a)
```
#pragma omp parallel
{
    for (int x = 0; x < len; x++){
        *(A+x) = x;
    }
}
```

Is the code:

A) Always Incorrect

B) Sometimes Incorrect

C) Always Correct, Slower than Serial

D) Always Correct, Speed relative to Serial depends on Caching Scheme

E) Always Correct, Faster than Serial